

# task1

December 4, 2020

## 1 Class Challenge: Image Classification of COVID-19 X-rays

### 2 Task 1 [Total points: 30]

#### 2.1 Setup

- This assignment involves the following packages: 'matplotlib', 'numpy', and 'sklearn'.
- If you are using conda, use the following commands to install the above packages:

```
conda install matplotlib
conda install numpy
conda install -c anaconda scikit-learn
```

- If you are using pip, use the following commands to install the above packages:

```
pip install matplotlib
pip install numpy
pip install sklearn
```

#### 2.2 Data

Please download the data using the following link: [COVID-19](#).

- After downloading 'Covid\_Data\_GradientCrescent.zip', unzip the file and you should see the following data structure:

```
|--all |--train |--test |--two |--train |--test
```

- Put the 'all' folder, the 'two' folder and this python notebook in the **same directory** so that the following code can correctly locate the data.

#### 2.3 [20 points] Binary Classification: COVID-19 vs. Normal

```
[1]: import os

import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator

os.environ['OMP_NUM_THREADS'] = '1'
```

```
os.environ['CUDA_VISIBLE_DEVICES'] = '-1'
tf.__version__
```

```
[1]: '2.3.1'
```

### Load Image Data

```
[2]: DATA_LIST = os.listdir('two/train')
DATASET_PATH = 'two/train'
TEST_DIR = 'two/test'
IMAGE_SIZE = (224, 224)
NUM_CLASSES = len(DATA_LIST)
BATCH_SIZE = 10 # try reducing batch size or freeze more layers if your GPU
↳ runs out of memory
NUM_EPOCHS = 40
LEARNING_RATE = 0.0005 # start off with high rate first 0.001 and experiment
↳ with reducing it gradually
```

### Generate Training and Validation Batches

```
[3]: train_datagen = ImageDataGenerator(rescale=1./
↳ 255, rotation_range=50, featurewise_center = True,
featurewise_std_normalization =
↳ True, width_shift_range=0.2,
height_shift_range=0.2, shear_range=0.
↳ 25, zoom_range=0.1,
zca_whitening = True, channel_shift_range =
↳ 20,
horizontal_flip = True, vertical_flip = True,
validation_split = 0.2, fill_mode='constant')

train_batches = train_datagen.
↳ flow_from_directory(DATASET_PATH, target_size=IMAGE_SIZE,
↳ shuffle=True, batch_size=BATCH_SIZE,
subset = "training", seed=42,
class_mode="binary")

valid_batches = train_datagen.
↳ flow_from_directory(DATASET_PATH, target_size=IMAGE_SIZE,
↳ shuffle=True, batch_size=BATCH_SIZE,
subset = "validation", seed=42,
class_mode="binary")
```

/Users/bellarocha/anaconda3/envs/tf2/lib/python3.7/site-packages/keras\_preprocessing/image/image\_data\_generator.py:342: UserWarning:

This ImageDataGenerator specifies `zca\_whitening` which overrides setting of `featurewise\_std\_normalization`.

```
warnings.warn('This ImageDataGenerator specifies '
```

Found 104 images belonging to 2 classes.

Found 26 images belonging to 2 classes.

**[10 points] Build Model** Hint: Starting from a pre-trained model typically helps performance on a new task, e.g. starting with weights obtained by training on ImageNet.

```
[4]: vgg16 = tf.keras.applications.VGG16(  
    include_top=False,  
    weights="imagenet",  
    input_shape= (224, 224, 3)  
)  
  
model = tf.keras.models.Sequential([  
    vgg16,  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(256, activation = 'relu', name = 'dense_feature'),  
    tf.keras.layers.Dropout(0.2),  
    tf.keras.layers.Dense(1, activation = 'sigmoid')  
)  
  
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 7, 7, 512)	14714688
flatten (Flatten)	(None, 25088)	0
dense_feature (Dense)	(None, 256)	6422784
dropout (Dropout)	(None, 256)	0
dense (Dense)	(None, 1)	257

Total params: 21,137,729  
Trainable params: 21,137,729  
Non-trainable params: 0

**[5 points] Train Model**

```
[5]: #FIT MODEL
print(len(train_batches))
print(len(valid_batches))

STEP_SIZE_TRAIN = train_batches.n//train_batches.batch_size
STEP_SIZE_VALID = valid_batches.n//valid_batches.batch_size

OPTIMIZER = tf.keras.optimizers.SGD(learning_rate = LEARNING_RATE)

model.compile(optimizer = OPTIMIZER, loss = tf.keras.losses.
↳BinaryCrossentropy(), metrics = ['accuracy'])

history = model.fit(train_batches, epochs = NUM_EPOCHS, steps_per_epoch =
↳STEP_SIZE_TRAIN, validation_data = valid_batches, validation_steps =
↳STEP_SIZE_VALID)

11
3

/Users/bellarocha/anaconda3/envs/tf2/lib/python3.7/site-
packages/keras_preprocessing/image/image_data_generator.py:720: UserWarning:
This ImageDataGenerator specifies `featurewise_center`, but it hasn't been fit
on any training data. Fit it first by calling `.fit(numpy_data)`.
  warnings.warn('This ImageDataGenerator specifies '
/Users/bellarocha/anaconda3/envs/tf2/lib/python3.7/site-
packages/keras_preprocessing/image/image_data_generator.py:739: UserWarning:
This ImageDataGenerator specifies `zca_whitening`, but it hasn't been fit on any
training data. Fit it first by calling `.fit(numpy_data)`.
  warnings.warn('This ImageDataGenerator specifies '

Epoch 1/40
10/10 [=====] - 75s 7s/step - loss: 0.7663 - accuracy:
0.4574 - val_loss: 0.6746 - val_accuracy: 0.5500
Epoch 2/40
10/10 [=====] - 69s 7s/step - loss: 0.6302 - accuracy:
0.6596 - val_loss: 0.5159 - val_accuracy: 0.9500
Epoch 3/40
10/10 [=====] - 73s 7s/step - loss: 0.6430 - accuracy:
0.5700 - val_loss: 0.5000 - val_accuracy: 0.9000
Epoch 4/40
10/10 [=====] - 80s 8s/step - loss: 0.6259 - accuracy:
0.6383 - val_loss: 0.5190 - val_accuracy: 0.6500
Epoch 5/40
10/10 [=====] - 72s 7s/step - loss: 0.5463 - accuracy:
0.7553 - val_loss: 0.4607 - val_accuracy: 0.8500
Epoch 6/40
10/10 [=====] - 75s 8s/step - loss: 0.4358 - accuracy:
0.7872 - val_loss: 0.4379 - val_accuracy: 0.9000
```

Epoch 7/40  
10/10 [=====] - 75s 7s/step - loss: 0.4806 - accuracy: 0.7660 - val\_loss: 0.2739 - val\_accuracy: 1.0000

Epoch 8/40  
10/10 [=====] - 76s 8s/step - loss: 0.4418 - accuracy: 0.8191 - val\_loss: 0.2376 - val\_accuracy: 1.0000

Epoch 9/40  
10/10 [=====] - 76s 8s/step - loss: 0.4260 - accuracy: 0.7660 - val\_loss: 0.2071 - val\_accuracy: 0.9500

Epoch 10/40  
10/10 [=====] - 77s 8s/step - loss: 0.3542 - accuracy: 0.8404 - val\_loss: 0.2407 - val\_accuracy: 1.0000

Epoch 11/40  
10/10 [=====] - 79s 8s/step - loss: 0.2748 - accuracy: 0.8830 - val\_loss: 0.1238 - val\_accuracy: 1.0000

Epoch 12/40  
10/10 [=====] - 73s 7s/step - loss: 0.2845 - accuracy: 0.9043 - val\_loss: 0.1411 - val\_accuracy: 1.0000

Epoch 13/40  
10/10 [=====] - 73s 7s/step - loss: 0.3900 - accuracy: 0.8298 - val\_loss: 0.1437 - val\_accuracy: 0.9500

Epoch 14/40  
10/10 [=====] - 73s 7s/step - loss: 0.2416 - accuracy: 0.9043 - val\_loss: 0.1467 - val\_accuracy: 0.9500

Epoch 15/40  
10/10 [=====] - 65s 7s/step - loss: 0.2489 - accuracy: 0.8830 - val\_loss: 0.1153 - val\_accuracy: 1.0000

Epoch 16/40  
10/10 [=====] - 65s 6s/step - loss: 0.2470 - accuracy: 0.8830 - val\_loss: 0.0926 - val\_accuracy: 0.9500

Epoch 17/40  
10/10 [=====] - 78s 8s/step - loss: 0.2997 - accuracy: 0.8500 - val\_loss: 0.0941 - val\_accuracy: 1.0000

Epoch 18/40  
10/10 [=====] - 65s 6s/step - loss: 0.1868 - accuracy: 0.9468 - val\_loss: 0.0926 - val\_accuracy: 1.0000

Epoch 19/40  
10/10 [=====] - 65s 7s/step - loss: 0.3663 - accuracy: 0.8298 - val\_loss: 0.1077 - val\_accuracy: 0.9500

Epoch 20/40  
10/10 [=====] - 65s 6s/step - loss: 0.2334 - accuracy: 0.8936 - val\_loss: 0.0976 - val\_accuracy: 0.9500

Epoch 21/40  
10/10 [=====] - 70s 7s/step - loss: 0.1820 - accuracy: 0.9149 - val\_loss: 0.1140 - val\_accuracy: 1.0000

Epoch 22/40  
10/10 [=====] - 69s 7s/step - loss: 0.1869 - accuracy: 0.9255 - val\_loss: 0.1229 - val\_accuracy: 0.9500

Epoch 23/40  
10/10 [=====] - 65s 7s/step - loss: 0.1670 - accuracy: 0.9149 - val\_loss: 0.0460 - val\_accuracy: 1.0000

Epoch 24/40  
10/10 [=====] - 69s 7s/step - loss: 0.1556 - accuracy: 0.9500 - val\_loss: 0.0847 - val\_accuracy: 0.9500

Epoch 25/40  
10/10 [=====] - 65s 6s/step - loss: 0.2279 - accuracy: 0.9043 - val\_loss: 0.1085 - val\_accuracy: 0.9500

Epoch 26/40  
10/10 [=====] - 65s 7s/step - loss: 0.1304 - accuracy: 0.9468 - val\_loss: 0.0629 - val\_accuracy: 0.9500

Epoch 27/40  
10/10 [=====] - 65s 6s/step - loss: 0.1960 - accuracy: 0.9043 - val\_loss: 0.1150 - val\_accuracy: 0.9500

Epoch 28/40  
10/10 [=====] - 65s 6s/step - loss: 0.2079 - accuracy: 0.9255 - val\_loss: 0.0500 - val\_accuracy: 1.0000

Epoch 29/40  
10/10 [=====] - 69s 7s/step - loss: 0.1393 - accuracy: 0.9400 - val\_loss: 0.1353 - val\_accuracy: 0.9000

Epoch 30/40  
10/10 [=====] - 65s 7s/step - loss: 0.1785 - accuracy: 0.9255 - val\_loss: 0.0658 - val\_accuracy: 1.0000

Epoch 31/40  
10/10 [=====] - 65s 7s/step - loss: 0.1428 - accuracy: 0.9362 - val\_loss: 0.0562 - val\_accuracy: 1.0000

Epoch 32/40  
10/10 [=====] - 69s 7s/step - loss: 0.1050 - accuracy: 0.9574 - val\_loss: 0.0490 - val\_accuracy: 1.0000

Epoch 33/40  
10/10 [=====] - 68s 7s/step - loss: 0.1268 - accuracy: 0.9574 - val\_loss: 0.0828 - val\_accuracy: 0.9500

Epoch 34/40  
10/10 [=====] - 65s 6s/step - loss: 0.1082 - accuracy: 0.9574 - val\_loss: 0.0233 - val\_accuracy: 1.0000

Epoch 35/40  
10/10 [=====] - 65s 6s/step - loss: 0.1178 - accuracy: 0.9574 - val\_loss: 0.0404 - val\_accuracy: 1.0000

Epoch 36/40  
10/10 [=====] - 64s 6s/step - loss: 0.0662 - accuracy: 0.9894 - val\_loss: 0.0327 - val\_accuracy: 1.0000

Epoch 37/40  
10/10 [=====] - 69s 7s/step - loss: 0.0703 - accuracy: 0.9900 - val\_loss: 0.0570 - val\_accuracy: 0.9500

Epoch 38/40  
10/10 [=====] - 65s 6s/step - loss: 0.1673 - accuracy: 0.9149 - val\_loss: 0.1675 - val\_accuracy: 0.9500



```
Epoch 39/40
10/10 [=====] - 65s 6s/step - loss: 0.2361 - accuracy:
0.9043 - val_loss: 0.0356 - val_accuracy: 1.0000
Epoch 40/40
10/10 [=====] - 69s 7s/step - loss: 0.0806 - accuracy:
0.9600 - val_loss: 0.0440 - val_accuracy: 1.0000
```

#### [5 points] Plot Accuracy and Loss During Training

```
[9]: import matplotlib.pyplot as plt

plt.plot(history.history['accuracy'], label='Train_acc')
plt.plot(history.history['val_accuracy'], label = 'Test_acc')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.6, 1.02])
plt.legend(loc='lower right')
plt.title('Accuracy over 40 Epochs')
plt.show()

plt.plot(history.history['loss'], label='Train_loss')
plt.plot(history.history['val_loss'], label = 'Test_loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.ylim([0, 0.7])
plt.legend(loc='upper right')
plt.title('Loss over 40 Epochs')
plt.show()
```





## Plot Test Results

```
[10]: import matplotlib.image as mpimg

test_datagen = ImageDataGenerator(rescale=1. / 255)
eval_generator = test_datagen.
    ↳flow_from_directory(TEST_DIR,target_size=IMAGE_SIZE,

    ↳batch_size=1,shuffle=True,seed=42,class_mode="binary")
eval_generator.reset()
pred = model.predict_generator(eval_generator,18,verbose=1)
for index, probability in enumerate(pred):
    image_path = TEST_DIR + "/" +eval_generator_filenames[index]
    image = mpimg.imread(image_path)
    if image.ndim < 3:
        image = np.reshape(image,(image.shape[0],image.shape[1],1))
        image = np.concatenate([image, image, image], 2)
    #         print(image.shape)


    pixels = np.array(image)
    plt.imshow(pixels)

    print(eval_generator_filenames[index])
    if probability > 0.5:
        plt.title("%.2f" % (probability[0]*100) + "% Normal")
    else:
        plt.title("%.2f" % ((1-probability[0])*100) + "% COVID19 Pneumonia")
    plt.show()
```


Found 18 images belonging to 2 classes.

18/18 [=====] - 4s 207ms/step


covid/nejmoa2001191\_f3-PA.jpeg




covid/nejmoa2001191\_f4.jpeg




covid/nejmoa2001191\_f5-PA.jpeg




covid/radiol.2020200490.fig3.jpeg




covid/ryct.2020200028.fig1a.jpeg




covid/ryct.2020200034.fig2.jpeg




covid/ryct.2020200034.fig5-day0.jpeg




covid/ryct.2020200034.fig5-day4.jpeg




covid/ryct.2020200034.fig5-day7.jpeg




normal/NORMAL2-IM-1385-0001.jpeg




normal/NORMAL2-IM-1396-0001.jpeg



normal/NORMAL2-IM-1400-0001.jpeg




normal/NORMAL2-IM-1401-0001.jpeg







normal/NORMAL2-IM-1406-0001.jpeg




normal/NORMAL2-IM-1412-0001.jpeg




normal/NORMAL2-IM-1419-0001.jpeg



normal/NORMAL2-IM-1422-0001.jpeg



normal/NORMAL2-IM-1423-0001.jpeg



## 2.4 [10 points] TSNE Plot

t-Distributed Stochastic Neighbor Embedding (t-SNE) is a widely used technique for dimensionality reduction that is particularly well suited for the visualization of high-dimensional datasets. After training is complete, extract features from a specific deep layer of your choice, use t-SNE to reduce the dimensionality of your extracted features to 2 dimensions and plot the resulting 2D features.

```
[18]: from sklearn.manifold import TSNE

intermediate_layer_model = tf.keras.models.Model(inputs=model.input,
                                                  outputs=model.
                                                  ↳get_layer('dense_feature').output)
tsne_data_generator = test_datagen.
↳flow_from_directory(DATASET_PATH,target_size=IMAGE_SIZE,
↳batch_size=1,shuffle=False,seed=42,class_mode="binary")

intermediate_layer = intermediate_layer_model.predict(tsne_data_generator)

intermediate_layer_TSNE = TSNE().fit_transform(intermediate_layer)

classes = tsne_data_generator.classes

colors = []

for i in range(len(classes)):
    if classes[i] == 0:
        colors.append('red')
    else:
        colors.append('blue')

plt.scatter(intermediate_layer_TSNE[:, 0], intermediate_layer_TSNE[:, 1], color=
↳colors)
plt.scatter(intermediate_layer_TSNE[:, 0][0], intermediate_layer_TSNE[:, 1][0],
↳color = colors[0], label = 'COVID-19')
plt.scatter(intermediate_layer_TSNE[:, 0][70], intermediate_layer_TSNE[:,
↳1][70], color = colors[70], label = 'Normal')
plt.legend()
plt.show()
```

Found 130 images belonging to 2 classes.

